




Intro to the Intel x64

Part 2

1




The Intel x64

It was simple at first...

2

The Intel x64

- The Intel x64 is the main processor used by servers, laptops, and desktops
- It has evolved continuously over a 40+ year period



Source: 2022


Source: 2022

3

3

The Original x86

- First "x86" was the 8086
- Released in 1978
- Attributes:
 - 16-bit registers
 - 16 registers
 - could access of 1MB of RAM (in 64KB blocks using a special "segment" register)



Source: 2022


Source: 2022

4

4

What to call the processor

- The classic term "x86" refers to the 32-bit and 16-bit processor family
- With move to 64-bit, the term "x64" is used to differentiate the newest design from the previous



Source: 2022

Source: 2022

5

5



Original x86 Registers

It was simple at first...

6

Original x86 Registers

- The original x86 contained 16 registers
- 8 can be used by your programs
- The other 8 are used for memory management



Summer 2022

Segmented State - Cook - CSU 35

7

7

Original x86 Registers

- The x86 processor has evolved continuously over the last 4 decades
- It first jumped to 32-bit, and then, again, to 64-bit
- The result is many of the registers have strange names

Summer 2022

Segmented State - Cook - CSU 35

8

8

Original x86 Registers

- 8 Registers can be used by your programs
 - Four General Purpose: **AX, BX, CX, DX**
 - Four pointer index: **SI, DI, BP, SP**
- The remaining 8 are restricted
 - Six segment: CS, DS, ES, FS, GS, SS
 - One instruction pointer: IP
 - One status register – used in computations

Summer 2022


Segmented State - Cook - CSU 35

9

9

Original General-Purpose Registers

- However, back then (and now too) it is very useful to store 8-bit values
- So, Intel chopped 4 of the registers in half
- These registers have generic names of A, B, C, D



Summer 2022


Segmented State - Cook - CSU 35

10

10

Original General-Purpose Registers

- The first and second byte can be used separately or used together
- Naming convention
 - high byte has the suffix "H"
 - low byte has the suffix "L"
 - for both bytes, the suffix is "X"



Summer 2022


Segmented State - Cook - CSU 35

11

11

Original General-Purpose Registers

- This essentially doubled the number of registers
- So, there are:
 - four 16-bit registers or
 - eight 8-bit registers
 - ...and any combination you can think off




Summer 2022

Segmented State - Cook - CSU 35

12

12

Last the 4 Registers



- The remaining 4 registers were not cut in half
- Used for storing indexes (for arrays) and pointers
- Their purpose
 - DI – destination index
 - SI – source index
 - BP – base pointer
 - SP – stack pointer


Summer 2022

Segmented State - Cook - CSU 35

13

13

Original 16-Bit Registers



Summer 2022

Segmented State - Cook - CSU 35

14

14

Evolution to 64-Bit Registers



This is going to hurt...

Summer 2022

15

Evolution to 32-bit

- When the x86 moved to 32-bit era, Intel expanded the registers to 32-bit
 - the 16-bit ones still exist
 - they have the prefix "e" for extended
- New instructions were added to use them




Summer 2022

Segmented State - Cook - CSU 35

16

16

Original Registers




Summer 2022

Segmented State - Cook - CSU 35

17

17

Expansion to 32-bit




Summer 2022

Segmented State - Cook - CSU 35

18


18

Original Registers



19

Expansion to 32-bit



20

Evolution to 64-bit

- At this point, Intel had decided to abandon the x86 in lieu of their new Itanium Processor
- The Itanium was a radically different design and was completely incompatible
- Advanced Micro Devices (AMD), to Intel's chagrin, decided to — *once again* — extend the x86



21

Evolution to 64-bit

- Registers were extended again
 - 64-bit registers have the prefix "*r*"
 - 8 additional registers were added
 - also, it is now possible to get 8-bit values from all registers (hardware is more consistent!)
- Some old, archaic, features were dropped



22


Evolution to 64-bit

- The AMD-64 was a huge commercial *success*
- The Itanium was a commercial *failure*
- Intel, dropped the Itanium and started making 64-bit x86 using AMD's design




23

Expansion to 64-bit



24

Expansion to 64-bit




Summer 2022

Background: State - Cook - CSU 28

25

25

Expansion to 64-bit




Summer 2022

Background: State - Cook - CSU 31

26

26

Expansion to 64-bit




Summer 2022

Background: State - Cook - CSU 32

27

27

New 64-bit Registers: R8...R15



Summer 2022

Background: State - Cook - CSU 33

28

28

64-Bit Register Table

Register	32 bit	16-bit	8-bit High	8-bit Low
rax	eax	ax	ah	al
rbx	ebx	bx	bh	bl
rcx	ecx	cx	ch	cl
rdx	edx	dx	dh	dl
rsi	esi	si		sil
rdi	edi	di		dil
rbp	ebp	bp		bpl
rsp	esp	sp		spl

Summer 2022

Background: State - Cook - CSU 34

29

29

64-Bit Register Table

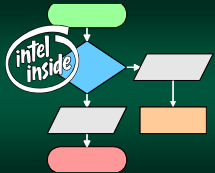
Register	32 bit	16-bit	8-bit High	8-bit Low
r8	r8d	r8w		r8b
r9	r9d	r9w		r9b
r10	r10d	r10w		r10b
r11	r11d	r11w		r11b
r12	r12d	r12w		r12b
r13	r13d	r13w		r13b
r14	r14d	r14w		r14b
r15	r15d	r15w		r15b

Summer 2022

Background: State - Cook - CSU 35

30

30



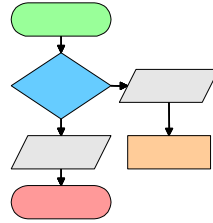
Basic Intel x86 Instructions

Feel the pow-wah of the x86!

31

Basic Intel x86 Instructions

- Each x86 instruction can have up to 2 operands
- Operands in x86 instructions are very versatile
- Each operand can be either a memory address, register or an immediate value



32

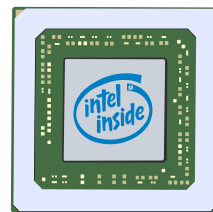
Types of Operands

- Registers
- Address in memory
- Register pointing to a memory address
- Immediate

33

Intel x86 Instruction Limits

- There are some limitations...
- Some instructions must use an immediate
- Some instructions require a specific register to perform calculations



34

Intel x86 Instruction Limits

- A register must always be involved
 - processors use registers for all activity
 - both operands cannot access memory at the same time
 - the processor has to have it at some point!*
- Also, obviously, the receiving field cannot be an immediate value

35

Instruction: Move

- The Intel Move Instruction combines transfer, load and store instructions under one name
- ... well, that's something the assembler does for us – but, we'll cover that soon
- "Move" is a tad confusing – it copies data

36

Instruction: Move

```
MOV destination, source
```

Immediate, Register, Memory

Register, Memory

37

Example: Move immediate

```
MOV rax, 42
```

Source is a immediate constant

Same as Java
`rax = 42;`

Destination is rax

38

Example: Transfer

```
MOV rbx, rax
```

Source is rax

Same as Java
`rbx = rax;`

Destination is rbx

39

Example: Load

```
MOV rax, total
```

"total" is memory location

Destination is rax

40

Example: Store

```
MOV counter, rax
```

Source is rax

Memory location named 'Counter'

41

Example: "A" Register

So many options!

```
mov al, 42      #low byte
```

```
mov ah, 13      #high byte
```

```
mov ax, 1947    #both bytes
```

42

Instruction: Add & Subtract

- The Add and Subtract instructions take two operands and store the result in the first operand
- This is the same as the **+=** and **-=** operators used in Visual Basic .NET, C, C++, Java, etc...



Summer 2022

Backwards Step - Cook - CS0 35

43

43

Instruction: Add

ADD *target* , *value*

Immediate, Register, Memory

Register, Memory

Summer 2022

Backwards Step - Cook - CS0 35

44

44

Example: Move register to memory

Move memory into rax

```
MOV rax, counter
ADD rax, 2
```

Same as Java
`rax += 2;`

Summer 2022


Backwards Step - Cook - CS0 35

45

45

Instruction: And & Or

- The Bitwise And & Bitwise Or instructions take two operands and stores the result in the second operand
- This is the same as the **&=** and **|=** operators used in C, C++, Java, etc...



Summer 2022

Backwards Step - Cook - CS0 35

46

46

Instruction: Logical And

AND *target* , *value*

Immediate, Register, Memory

Register, Memory

Summer 2022

Backwards Step - Cook - CS0 35

47

47

Example: Logical Or

```
#Convert 5 to ASCII '5'
MOV rax, 5
OR rax, 0x30
```

0011 0000

Summer 2022

Backwards Step - Cook - CS0 35

48

48

Call Instruction

- The *Call Instruction* causes the processor to start running instructions at a specified memory location (a subroutine)
- Subroutines are analogous to the functions you wrote in Java
- Once it completes, execution returns from the subroutine and continues after the call

Summer 2022

Secrets@Stan - CS61B

49

49

Call Instruction

CALL *address*

Usually a label – a constant that holds an address

Summer 2022

Secrets@Stan - CS61B

50

50

Example: Print an integer

#Using the CSC35 library

MOV rdx, 1846

CALL PrintInt

This name is an address

Summer 2022

Secrets@Stan - CS61B

51

51