



## Memory

Part 4

1



## von Neumann Architecture

The Information Superhighway

2

### von Neumann Machine Architecture

- Modern computers are based on the design of John von Neumann
- His design greatly simplified the construction of (and use) computers



Fall 2022

Seacornville State - Oak - CSU 25

3

3

### Some von Neumann Attributes

1. Programs are stored and executed in memory
2. Separation of processing from memory
3. Different system components communicate over a shared bus



Fall 2022

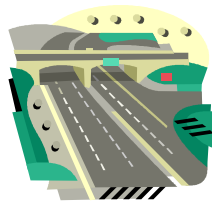
Seacornville State - Oak - CSU 25

4

4

### The Bus

- Electronic pathway that transports data between components
- Think of it as a "highway"
  - data moves on shared paths
  - otherwise, the computer would be very complex



Fall 2022

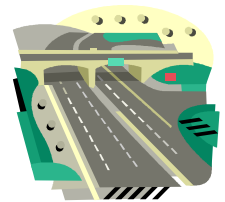
Seacornville State - Oak - CSU 25

5

5

### System Bus

- The information sent on the memory bus falls into **3** categories
- Three sets of signals
  - address bus
  - data bus
  - control bus



Fall 2022

Seacornville State - Oak - CSU 25

6

6

## Address Bus

- Used by the processor to access a specific piece of data
- This "address" can be
  - a specific byte in memory
  - unique IO port
  - etc...
- The more bits it has, the more memory can be accessed



Fall 2022

Secretariat State - Cook - CSU 35

7

7

## Address Bus Size Examples

- 8-bit  $\rightarrow 2^8 = 256$  bytes
- 16-bit  $\rightarrow 2^{16} = 64$  KB (65,536 bytes)
- 32-bit  $\rightarrow 2^{32} = 4$  GB (4,294,967,296 bytes)
- 64-bit  $\rightarrow 2^{64} = 18$  EB (18,446,744,073,709,551,616)



Fall 2022

Secretariat State - Cook - CSU 35

8

8

## Historic Address Sizes

- Intel 8086
  - original 1982 IBM PC
  - 20-bit address bus (1 MB)
  - only 640 KB usable for programs
- MOS 6502 computers
  - Commodore 64, Apple II, Nintendo, etc...
  - 16-bit address bus (64 KB)

Fall 2022

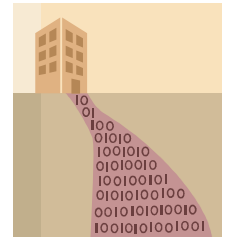
Secretariat State - Cook - CSU 35

9

9

## Data Bus

- The actual data travels over the *data bus*
- The number of bits that the processor uses – as its natural unit of data – is called a *word*



Fall 2022

Secretariat State - Cook - CSU 35

10

10

## Data Bus

- Typically we define a system by word size
- Example:
  - 8-bit system uses 8 bit words
  - 16-bit system uses 16 bits (2 bytes) words
  - 32-bit system uses 32 bits (4 bytes) words
  - etc...

Fall 2022

Secretariat State - Cook - CSU 35

11

11

## Control Bus

- The *control bus* controls the timing and synchronizes the subsystems
- Specifies what is happening
  - read data
  - write data
  - reset
  - etc...

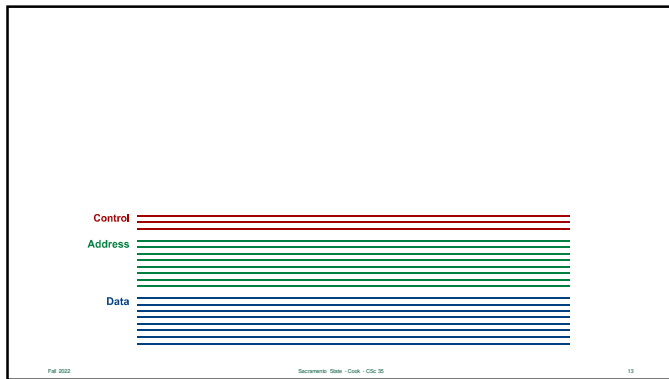


Fall 2022

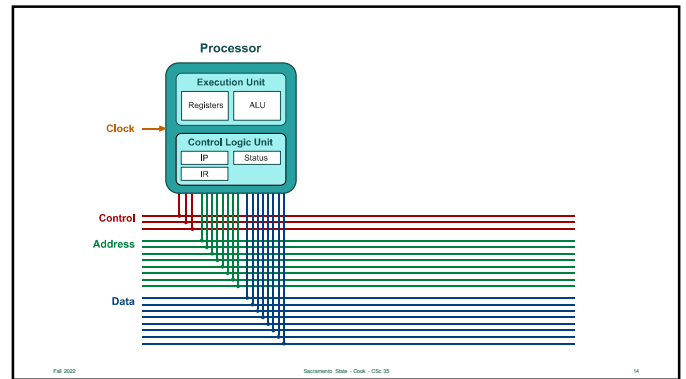
Secretariat State - Cook - CSU 35

12

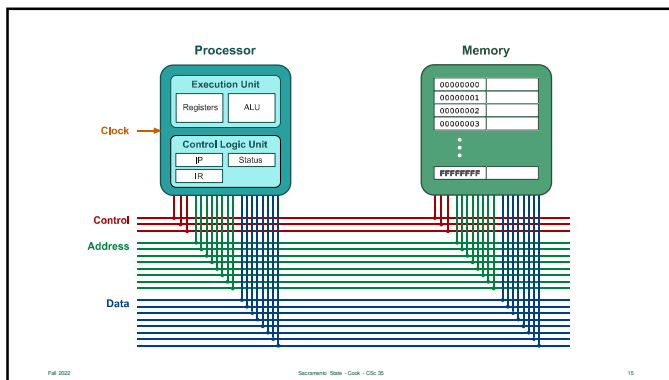
12



13



14



15

### von Neumann Architecture Today

- Because of the emphasis on memory, most real-world systems use a modified version of his design
- In particular, they have a special high-speed bus between the processor and memory




Diagram showing a processor and memory connected by three buses: Control (red), Address (green), and Data (blue). The processor is divided into an Execution Unit (Registers, ALU) and a Control Logic Unit (IP, Status, IR). The memory is shown as a green box with a list of addresses (00000000, 00000001, 00000002, 00000003, ..., FFFFFFFF). The buses are represented by horizontal lines. The diagram is labeled with 'Fall 2022' and 'Srinivasan, Srinivasan - CSU 35'.

16

### von Neumann Architecture Today

- Think of it as a diamond-lane on a freeway
- ... or as high-speed rail – which has a fixed source and destination and goes faster than the freeway




Diagram showing a processor and memory connected by three buses: Control (red), Address (green), and Data (blue). The processor is divided into an Execution Unit (Registers, ALU) and a Control Logic Unit (IP, Status, IR). The memory is shown as a green box with a list of addresses (00000000, 00000001, 00000002, 00000003, ..., FFFFFFFF). The buses are represented by horizontal lines. The diagram is labeled with 'Fall 2022' and 'Srinivasan, Srinivasan - CSU 35'.

17

### Accessing Data

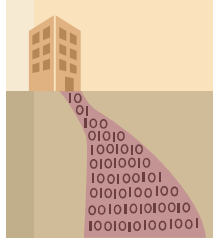


Diagram showing a processor and memory connected by three buses: Control (red), Address (green), and Data (blue). The processor is divided into an Execution Unit (Registers, ALU) and a Control Logic Unit (IP, Status, IR). The memory is shown as a green box with a list of addresses (00000000, 00000001, 00000002, 00000003, ..., FFFFFFFF). The buses are represented by horizontal lines. The diagram is labeled with 'Fall 2022' and 'Srinivasan, Srinivasan - CSU 35'.

18

## Accessing Data

- Processors use registers to hold data being computed
- So, how is data put into the registers to begin with?
- Data can come from two major sources



Fall 2022

Secrets: State - Cook - CSU 35

19

19

## Immediates

- In programming, it is common to assign a constant to a variable
- As you can imagine, this will also be quite common with instructions



Fall 2022

Secrets: State - Cook - CSU 35

20

20

## Immediates

- When a constant is stored as part of instruction, it is called an *immediate*
- Once the instruction is loaded by the processor, it is "immediately" available from the IR – hence, the name



Fall 2022

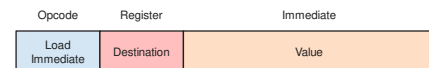
Secrets: State - Cook - CSU 35

21

21

## Load Immediate

- A *Load Immediate* instruction, stores a constant into a register
- The instruction must store the destination register and the immediate value



Fall 2022

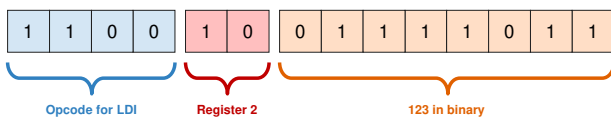
Secrets: State - Cook - CSU 35

22

22

## Load Immediate Example (not x86)

**LDI r2, 123**



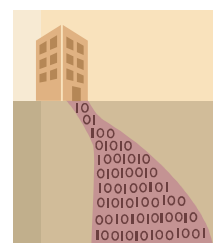
Fall 2022

Secrets: State - Cook - CSU 35

23

23

## Copying Data



- Processors have a number of instructions that can copy data
- Each has a unique name –*not surprising since each does something different*

Fall 2022

Secrets: State - Cook - CSU 35

24

24

## Transfer

- A *Transfer* instruction, copies the contents of one instruction into another
- The instruction must store both the destination and source register



Fall 2022

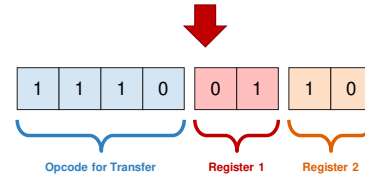
Sebastien Balle - CSE 331

25

25

## Transfer Example (not x86)

TRA r1, r2



Fall 2022

Sebastien Balle - CSE 331

26

26

## Effective Addresses

- Processors also have the ability to access memory
- Since memory is a massive array, the processor needs to know the address



Fall 2022

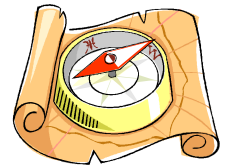
Sebastien Balle - CSE 331

27

27

## Effective Addresses

- The *effective address* is used to access memory
- Often it is created by combining multiple values
- ... but we will cover that later in the semester



Fall 2022

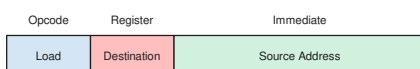
Sebastien Balle - CSE 331

28

28

## Load

- A *Load* instruction, reads data from memory (at a specified address)
- This data is then stored into the destination register



Fall 2022

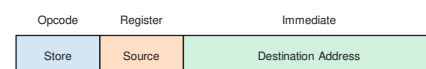
Sebastien Balle - CSE 331

29

29

## Store

- A *Store* instruction, writes data from a register into the specified address
- Note: the structure is identical to Load

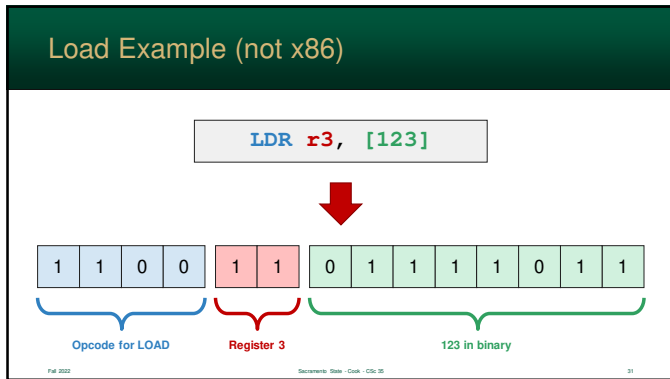


Fall 2022

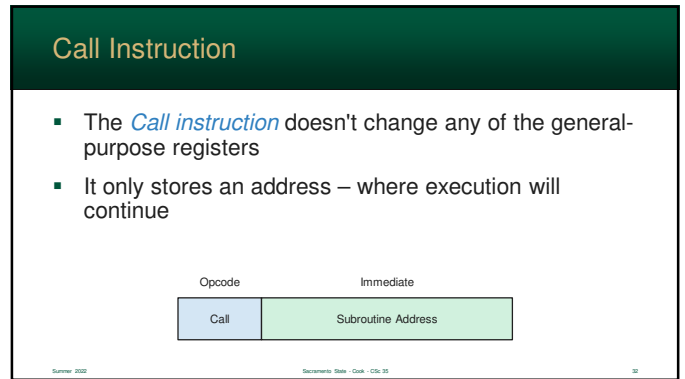
Sebastien Balle - CSE 331

30

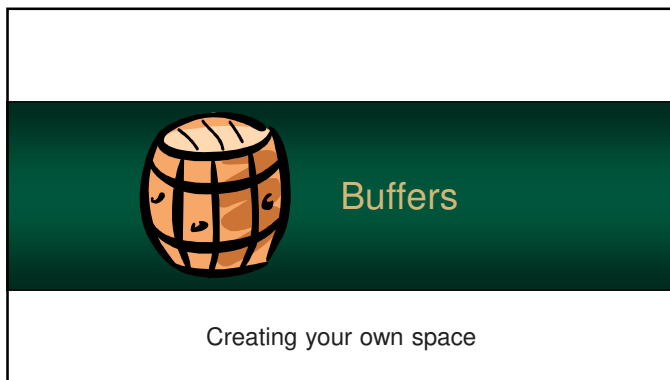
30



31



32



33

### Buffers

- A *buffer* is any allocated block of memory that contains data
- This can hold anything:
  - text
  - image
  - file
  - etc....

Fall 2022    Sacramento State - CSIS - CSIS 35    34

34

### Buffers

- There are several assembly *directives* which will allocate space
- We have covered a few of them, but there are many – all with a specific purpose

Fall 2022    Sacramento State - CSIS - CSIS 35    35

35

### A few directives that create space

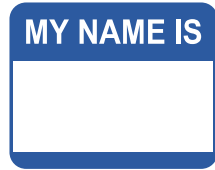
Directive	What it does
<code>.ascii</code>	Allocate enough space to store an ASCII string
<code>.quad</code>	Allocate 8-byte blocks with initial value(s)
<code>.byte</code>	Allocate byte(s) with initial value(s)
<code>.space</code>	Allocate any <i>size</i> of empty bytes (with initial values).

Fall 2022    Sacramento State - CSIS - CSIS 35    36

36

## Labels are addresses

- Labels are used to keep track of memory locations
- They are stored, by the assembler, in a table
- Whenever a label is used in the program, the assembler substitutes the address



Fall 2022

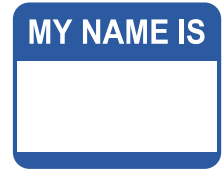
Segment: State - CSIS - CSIS 35

37

37

## Labels are addresses

- The table of labels is stored in the *object file*
- That way the linker can resolve any unknown labels
- After the program is linked into an executable, only addresses exist. No labels.



Fall 2022

Segment: State - CSIS - CSIS 35

38

38

## Quad Directive

Value:

.quad 74

Let's assume Value = 2000



2000	4A
2001	00
2002	00
2003	00
2004	00
2005	00
2006	00
2007	00

Fall 2022

Segment: State - CSIS - CSIS 35

39

39

## ASCII Directive Creates a Buffer

Text:

.ascii "Hello\0"

This label will store an address... once the assembler finds where to store it.

Creates 6 bytes to store Hello. They are stored consecutively.

Fall 2022

Segment: State - CSIS - CSIS 35

40

40

## Bytes are stored consecutively

Text:

.ascii "Hello\0"

Let's assume Text = 2000



2000	48	H
2001	65	e
2002	6C	l
2003	6C	l
2004	6F	o
2005	00	\0

Fall 2022

Segment: State - CSIS - CSIS 35

41

41

## Same Thing!

Text:

.byte 'H'  
.byte 'e'  
.byte 'l'  
.byte 'l'  
.byte 'o'  
.byte 0

Created byte by byte

Null character. Marks the end of a string

Fall 2022

Segment: State - CSIS - CSIS 35

42

42

## This works too!

```
Text:
.ascii "Hello"
.byte 0
```

Directives just create space. So, this creates a byte after the ASCII text.

43

## Create a Buffer of Any Size

```
Text:
.space 30
```

Create 30 bytes (defaults to 0x20 which is a space)

44

## Create a Buffer of Any Size

```
Text:
.space 30, 0
```

Create 30 bytes. All of which are 0.

45



## Direct Addressing

Using memory... finally

46

## Direct Addressing

- In *direct addressing*, the processor reads data directly from an address
- Commonly used to:
  - get a value from a "variable"
  - read items in an array
  - etc...



47

## Direct Addressing



48



## Direct in Java

- The following, for comparison, is the equivalent in Java
- The memory, at the address total, is loaded into `rdx`

```
// rdx = Memory[total];  
mov rdx, total
```

49

## Example: Direct Load

```
.intel_syntax noprefix  
.data  
funds:  
.quad 100
```

64 bit integer  
with an initial value of 100.

```
.text  
.global _start  
_start:  
mov rdx, funds
```

Read 8 bytes at this address.  
Doesn't store the address in `rdx`.

50

## Example: Direct Store

```
.intel_syntax noprefix  
.data  
funds:  
.quad 100  
  
.text  
.global _start  
_start:  
mov rdx, 5000  
mov funds, rdx
```

Store `rdx` into Address "funds"

51

## Example: Direct Store 2

```
.intel_syntax noprefix  
.data  
funds:  
.quad 100  
  
.text  
.global _start  
_start:  
call ScanInt  
mov funds, rdx
```

You can store inputted values.

52

## Direct in Java

- **Note:** this a shortcut notation
- The full notation would use square brackets
- The assembler recognizes the difference automatically

```
// rdx = Memory[total];  
mov rdx, total
```

53

## Direct in Java

- You can use the square-brackets if you want
- This way it explicitly show *how* the label is being used – it's a matter of preference

```
// rdx = Memory[total];  
mov rdx, [total]
```

54

## Example: Direct

```
.intel_syntax noprefix
.data
funds:
    .quad 100

.text
.global _start
_start:
    mov rdx, [funds]
```

A bit more descriptive

55

## Load Effective Address

- Load Effective Address stores the actual address into a register
- It doesn't access memory

```
// rdx = total;
lea rdx, total
```

56

## When to use `mov` and `lea`

The difference is huge!

57

## When to use `mov` and `lea`

- Knowing when to use an address **or** the data *located at that address* is vital
- Using the wrong one can cause your program to malfunction or crash

58

## Cause of the Segmentation Fault

- This is one of the most common mistakes in assembly programming

59

## Using Move Correctly

```
.intel_syntax noprefix
.data
Year:
    .quad 1947

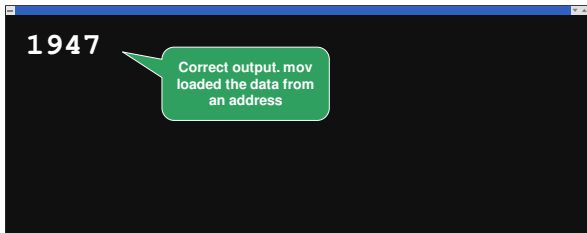
.text
.global _start
_start:
    mov rdx, Year
    call PrintInt
```

Creates 8 bytes

mov loads the data located at the address Year

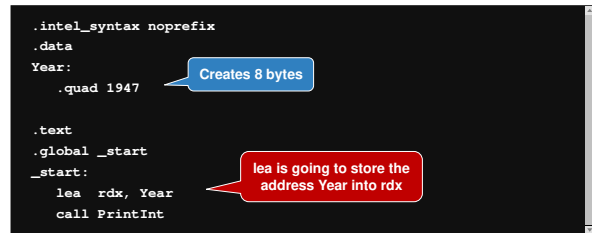
60

## Using move Correctly: Output



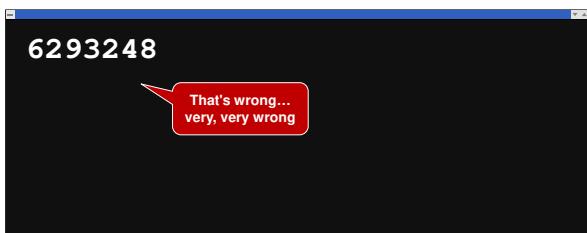
61

## Using `lea` by accident



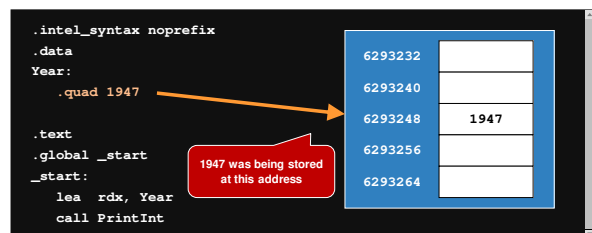
62

## Using `lea` by accident



63

## Why it Failed



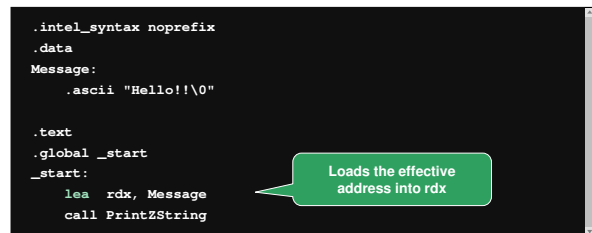
64

## Sometimes, You Need the Address

- Of course, sometimes, you do need an address
- For example, `PrintZString`
  - needs to know where the string is located so it can print a series of characters
  - so, it requires an address
  - `lea` is necessary

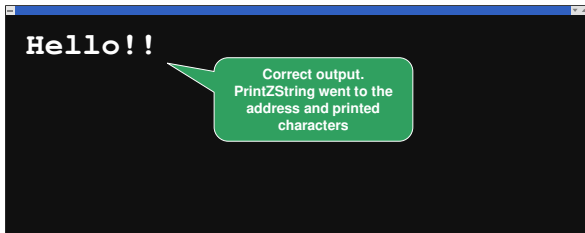
65

## Using `lea` correctly



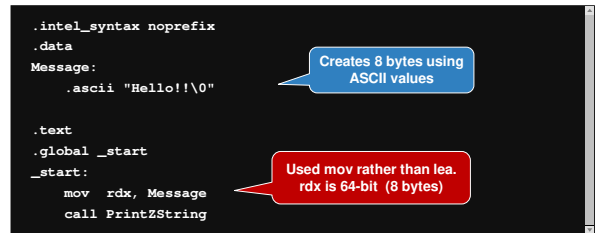
66

## Using `lea` correctly: Output



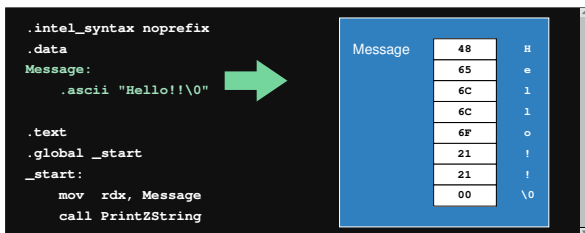
67

## Cause of the Segmentation Fault



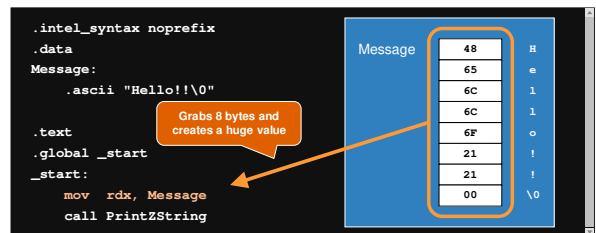
68

## Cause of the Segmentation Fault



69

## Cause of the Segmentation Fault



70

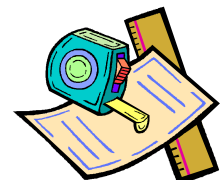
## Sizing Instructions

How many bytes are you using?

71

## Sizing Instructions

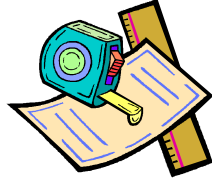
- The Intel can load/store 1-byte, 2-byte, 4-byte or 8-byte values
- The assembler knows (by *looking at the size of the register*) how much many bytes you want to load/store



72

## Sizing Instructions

- However, sometimes the number of bytes (1, 2, etc..) can't be determined
- In this case, the assembler will report an error
- ... since it doesn't know how to encode the instruction



Fall 2022

Section 10: Assembly - CS: 31

73

73

## Example: How Many Bytes?

```
.intel_syntax noprefix
.data
total:
    .quad 0

.text
.global _start
_start:
    mov total, 50
```

total is a target address. It doesn't have any implied size.

Fall 2022

Section 10: Assembly - CS: 31

74

74

## Example: How Many Bytes?

```
.intel_syntax noprefix
.data
total:
    .quad 0

.text
.global _start
_start:
    mov total, 50
```

How many bytes is this? The value 50 can be stored in 1, 2, 4, or 8 bytes.

Fall 2022

Section 10: Assembly - CS: 31

75

75

## How Many Bytes?

- If the assembler can't infer how many bytes to access, it'll report *"ambiguous operand size"*
- To address this issue...
  - GAS assembly allows you place a single character after the instruction's mnemonic
  - this suffix will tell the assembler how many bytes will be accessed during the operation

Fall 2022

Section 10: Assembly - CS: 31

76

76

## How Many Bytes

Suffix	Name	Size
b	byte	1 byte
s	short	2 bytes
l	long	4 bytes
q	quad	8 bytes

Fall 2022

Section 10: Assembly - CS: 31

77

77

## Example: Suffix Used

```
.intel_syntax noprefix
.data
total:
    .quad 0

.text
.global _start
_start:
    movq total, 50
```

Now the assembler knows you mean "move quad".


Fall 2022

Section 10: Assembly - CS: 31

78

78

# Endianness




The "proper" order of things

79

## So Many Bytes...

- On a 64-bit system, each word consists of 8 bytes
- So, when any 64-bit value is stored in memory, each of those 8 bytes must be stored
- However, question remains: *What order do we store them?*



80

## Example Unsigned Integer (4 Byte)

1,188,852,977

46	DC	74	F1
----	----	----	----

Most significant Byte (MSB)

Least significant Byte (LSB)

81

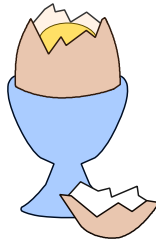
## So Many Bytes...

- Do we store the least-significant byte (LSB) first, or the most-significant (MSB)?
- As long as a system always follows the same format, then there are no problems
- ... but different system use different approaches

82

## Big Endian vs. Little Endian

- Big-Endian approach
  - store the MSB first
  - used by Motorola & PowerPC
- Little-Endian approach
  - store the LSB first
  - used by Intel



83

## Big Endian vs. Little Endian

46	DC	74	F1
----	----	----	----

Big Endian

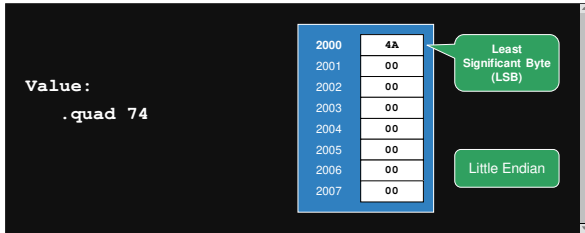
0	46
1	DC
2	74
3	F1

Little Endian

0	F1
1	74
2	DC
3	46

84

## Assuming Value is located at 2000



85

## No "End" to Problems

- *There is a problem...*  
if two systems use different formats, data will be interpreted incorrectly!
- If how the read differs from how it is stored, the data will be mangled



86

## No "End" to Problems

- For example:
  - a **little**-endian system reads a value stored in **big**-endian
  - a **big**-endian system reads a value stored in **little**-endian
- Programmers must be conscience of this whenever binary data is accessed



87

## No "End" to Problems

- So, whenever data is read from secondary storage, you cannot assume it will be in your processor's format
- This is compounded by file formats (gif, jpeg, mp3, etc...) which are also inconsistent



88

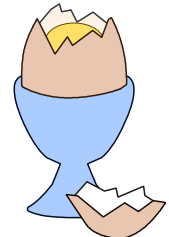
## Example File Format Endianness

File Format	Endianness
Adobe Photoshop	Big Endian
Windows Bitmap (.bmp)	Little Endian
GIF	Little Endian
JPEG	Big Endian
MP4	Big Endian
ZIP file	Little Endian

89

## So... who is correct?

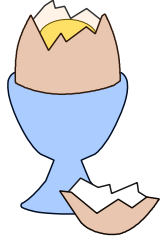
- So, what is the correct and superior format?
- Is it Intel (little endian)?
- ...or the PowerPC (big endian) correct?



90

## So... who is correct?

- In reality neither side is superior
- Both formats are equally correct
- Both have minor advantages in assembly... but nothing huge



Fall 2022

Sacramento State - Cook - CSU 38

91

91

## Gulliver's Travels



Fall 2022

Sacramento State - Cook - CSU 35

92

92